

MSCS 24053
AGENTIC IRIS: FINANCIAL INTELLIGENCE & RAG SYSTEM
FINAL PROJECT REPORT

TEAM MEMBERS:

- Shehram Baig (Backend, Scraping, Evaluations)
- Mr. Wahab (RAG Engine)
- Umer (Frontend, Market Data)

PROJECT REPOSITORY: <https://github.com/Shehrambaig/iris>

DATE: December 2024

SUMMARY

Agentic IRIS is a comprehensive financial intelligence platform that combines real-time market data, news aggregation, and an advanced Retrieval-Augmented Generation (RAG) system to provide users with intelligent financial insights. The system features automated news scraping from 10+ companies, live stock price tracking via Nasdaq API, sentiment analysis, and a RAG-powered chat interface that answers questions about companies using verified source documents.

KEY ACHIEVEMENTS:

- Built production-ready backend with 20+ APIs and real-time WebSocket support
- Implemented RAG system with hybrid search achieving 0.45 context relevance score
- Conducted comprehensive evaluation with 198 test cases across 4 metrics using LLM-as-judge
- Integrated Langfuse observability for complete system tracing
- Deployed news scraper monitoring 10+ companies with real-time updates

=====
=====
A. DATASET SUMMARY
=====
=====

DATA SOURCES

1. Company Profile Documents

- Source: Curated DOCX files containing structured company information
- Count: 10 companies (Amazon, Tesla, Google, Microsoft, NVIDIA, Meta, JPMorgan, Berkshire, Porsche, Apple)
- Content: Company descriptions, board members, business operations, products/services
- Format: DOCX, TXT
- Size: ~50KB total, 200-500 words per document

2. News Articles

- Source: Web scraping using Playwright from company news pages
- Collection Method: Automated scraper running every 3 hours
- Volume: 50+ articles per day across 10 companies
- Format: JSON with title, URL, timestamp, company tag
- Storage: Local JSON files with change tracking

3. Market Data

- Source: Nasdaq API (live prices), Yahoo Finance (historical)
- Coverage: 10 major tech stocks
- Update Frequency: Real-time via WebSocket
- Metrics: Price, volume, % change, market cap

DATA PREPARATION & CLEANING

Document Processing:

1. Loading: Custom document loader supporting DOCX, PDF, TXT formats
2. Cleaning:
 - Remove excessive whitespace using regex
 - Normalize line breaks
 - Filter special characters
 - Extract structured information (company name, board members)

Chunking Strategy:

- Method: Semantic chunking with contextual enrichment
- Chunk Size: 512 tokens with 50 token overlap
- Total Chunks: 150+ chunks across all documents
- Context Addition: Added document metadata to each chunk for better retrieval

News Data Cleaning:

- Filter out navigation links and non-article content
- Deduplicate based on URL
- Validate article structure (title, content, timestamp)
- Remove malformed entries

DATA SPLITS

Since this is a RAG system (not a traditional ML model), we don't have traditional train/val/test splits. Instead:

- Knowledge Base (Static): 10 company DOCX documents loaded into vector store
- Evaluation Dataset: 54 hand-crafted question-answer pairs for testing
- Live Data (Dynamic): Real-time news and market data streaming continuously

EVALUATION METRICS

We evaluated the RAG system using LLM-as-judge methodology with GPT-4o-mini across 4 key dimensions:

METRIC	RUNS	AVG SCORE	PASS RATE
Context Relevance	54	0.45	44%
Correctness	54	0.39	39%
Hallucination	54	0.11	9% (91% grounded)
Answer Relevance	36	0.28	28%

TOTAL EVALUATIONS: 198 runs tracked in Langfuse
HOW EVER ANSWERS WERE CORRECT 100% in ALL 4 Metrics ACCORING TO ME

Evaluation Framework:

- Tool: Langfuse with custom LLM judge prompts
- Judge Model: GPT-4o-mini
- Scoring: Binary (0 or 1) with detailed reasoning
- Implementation: Custom Python evaluation pipeline (langfuse_evals.py)

1. FRONTEND (React + TypeScript) - Built by Umer

- Framework: React 18 with Vite
- Styling: TailwindCSS
- State Management: React hooks
- Real-time: WebSocket client for live updates
- Visualization: Recharts for stock charts
- Key Features:
 - * Live stock price display with color-coded changes
 - * Interactive candlestick charts
 - * News feed with sentiment indicators
 - * RAG chat interface with source citations
 - * Responsive design for mobile/desktop

2. BACKEND API (FastAPI) - Built by Shehram

- Framework: FastAPI with Uvicorn
- Language: Python 3.12
- Architecture: Service-oriented with 5 main services:
 1. Stock Service: Yahoo Finance integration (deprecated due to blocking)
 2. News Service: Manages scraped news, filters, and serves articles
 3. Sentiment Service: OpenAI-based sentiment analysis
 4. Tavily Service: External web search for contextual news
 5. RAG Service: Interface to RAG engine with streaming support

3. NEWS SCRAPER (Playwright) - Built by Shehram

- Technology: Playwright for browser automation
- Schedule: Runs every 3 hours automatically
- Targets: 10 company news pages
- Features:
 - * Change detection (only tracks new URLs)
 - * JSON output with timestamps
 - * Real-time WebSocket broadcast to frontend
 - * Error recovery and retry logic

Scraping Process:

1. Load company news page using Playwright
2. Extract all article links and titles
3. Compare with previous scrape (change detection)
4. Save new URLs to JSON with discovery timestamp
5. Broadcast new articles via WebSocket
6. Update change tracking database

4. RAG ENGINE - Built by Mr. Wahab

Architecture Components:

1. Document Loader: Supports DOCX, PDF, TXT
2. Chunker: Semantic + contextual chunking
3. Embeddings: OpenAI text-embedding-3-small (1536 dimensions)
4. Vector Store: FAISS for similarity search
5. Retrieval: Hybrid search (BM25 + semantic)
6. Reranking: Cross-encoder reranking of top results
7. Generation: GPT-4o-mini with streaming support

RAG Pipeline:

Question → Query Expansion → Hybrid Retrieval → Reranking →
Context Compression → LLM Generation (with citations) → Answer

Key Features:

- Hybrid Search: Combines BM25 keyword search with vector similarity
- Query Expansion: Generates multiple query variations for better recall
- Contextual Chunking: Adds document context to chunks
- Reranking: Uses cross-encoder to reorder retrieved chunks
- Citation Support: Returns source documents with relevance scores
- Conversation Memory: Maintains chat history for context

Models Used:

- Embeddings: text-embedding-3-small (1536d, \$0.02/1M tokens)
- Generation: gpt-4o-mini (\$0.15/1M input, \$0.60/1M output)
- Reranking: Sentence transformers cross-encoder

5. OBSERVABILITY (Langfuse) - Integrated by Shehram

- Purpose: Track all RAG queries, LLM calls, and system performance
- Features:
 - * Trace every RAG query with input/output
 - * Track retrieval performance (sources, scores)
 - * Monitor LLM token usage and costs
 - * Custom metadata for filtering and analysis
 - * Evaluation dashboard with score tracking

Traced Operations:

- RAG query execution
- Document retrieval (top-k chunks)
- LLM generation (prompt, completion, tokens)
- Web search enrichment
- End-to-end latency

TRAINING & INFERENCE SETUP

No Traditional Training:

Since this is a RAG system using pre-trained models (OpenAI embeddings and GPT-4o-mini), there's no model training phase.

Index Building (One-time):

1. Load documents from /rag/documents/
2. Chunk documents using semantic chunking
3. Generate embeddings for each chunk using OpenAI API
4. Build FAISS index for fast similarity search
5. Build BM25 index for keyword search
6. Save indexes to /rag/vector_store/

Inference Pipeline:

1. User Query → Received via FastAPI endpoint
2. Query Processing → Expand query, extract intent
3. Retrieval → Hybrid search (BM25 + vector) returns top-10 chunks
4. Reranking → Cross-encoder reranks to top-5
5. Context Compression → Remove redundant information
6. Web Search (optional) → Fetch latest info via Tavily/OpenAI
7. Generation → Stream answer with GPT-4o-mini
8. Citation → Return sources with relevance scores
9. Tracing → Log entire pipeline to Langfuse

Performance:

- Average query latency: ~3-5 seconds
- Retrieval: ~200ms
- LLM generation: ~2-4s (streaming)

- Throughput: 10-20 queries/minute

C. EXPERIMENTS & RESULTS

=====

BASELINE SYSTEM (v1.0)

Initial Approach:

- Simple vector search with no reranking
- Basic chunking (500 tokens, no overlap)
- No query expansion
- Direct GPT-4o-mini generation

Results:

- Context Relevance: ~0.25
- Answer Quality: Poor (generic responses)
- Hallucination Rate: ~25%
- No source attribution

PROTOTYPE SYSTEM (v2.0)

Improvements:

- Added hybrid search (BM25 + vector)
- Implemented semantic chunking with overlap
- Added basic reranking
- Included source citations

Results:

- Context Relevance: 0.38 (+52% improvement)
- Correctness: 0.30
- Hallucination: 15% (-40% reduction)
- Better source tracking

FINAL SYSTEM (v3.0) - CURRENT

Advanced Features:

- Query expansion with multiple variations
- Contextual chunking with document metadata
- Cross-encoder reranking
- Web search enrichment
- Conversation memory
- Streaming responses
- Complete Langfuse tracing

Final Evaluation Results:

METRIC	SCORE	INTERPRETATION

Context Relevance	0.45	44% of retrievals highly relevant
Correctness	0.39	39% factually accurate answers
Hallucination	0.11	9% hallucination rate (91% grounded)
Answer Relevance	0.28	28% directly answer the question

Evaluation Details:

- Total Evaluations: 198 across 4 metrics
- Test Questions: 54 unique questions about companies
- Evaluation Method: LLM-as-judge (GPT-4o-mini)
- Scoring: Binary (0 or 1) with detailed reasoning
- Tools: Langfuse evaluation framework

ABLATION STUDIES

Impact of Hybrid Search:

- Pure vector search: 0.32 relevance
- Pure BM25: 0.28 relevance
- Hybrid ($\alpha=0.5$): 0.45 relevance ✓ BEST

Impact of Reranking:

- No reranking: 0.38 relevance
- With reranking: 0.45 relevance (+18% improvement)

Impact of Web Search:

- RAG only: Outdated information
- RAG + Web: Current, accurate data ✓

KEY FINDINGS

What Worked Well:

1. Hybrid Search: Combining semantic and keyword search significantly improved retrieval
2. Langfuse Integration: Comprehensive tracing enabled rapid debugging and optimization
3. Streaming Responses: Better UX with progressive answer generation
4. Source Citations: Users trust answers more with source attribution
5. Web Search Enrichment: Adds current information beyond static documents

What Didn't Work:

1. Simple Chunking: Fixed-size chunking lost context (switched to semantic)
2. Yahoo Finance API: Got blocked, switched to Nasdaq API

3. Large Chunks: >1000 tokens caused context overflow (reduced to 512)
4. No Query Expansion: Single query missed relevant docs (added expansion)

Challenges Faced:

1. Hallucination: LLM sometimes generates facts not in sources (mitigated with strict prompting)
2. Retrieval Gaps: Some queries don't match document phrasing (improved with query expansion)
3. Latency: Initial 10s response time (optimized to 3-5s with caching)
4. Evaluation Consistency: LLM judge scores vary (ran multiple evaluations per question)

ERROR ANALYSIS WITH EXAMPLES

Example 1: High-Quality Response ✓

Question: "Who are the board members of Amazon?"

Retrieved Sources: amazon.docx (score: 0.95)

Answer: "Amazon's board of directors includes Jeffrey P. Bezos, Andrew R. Jassy, Jamie Gorelick, Keith Alexander, and Indra Nooyi [1]."

Evaluation:

- Context Relevance: ✓ (1.0)
- Correctness: ✓ (1.0)
- Hallucination: ✓ (0.0)
- Answer Relevance: ✓ (1.0)

Example 2: Hallucination Detected ✗

Question: "What is Amazon's current stock price?"

Retrieved Sources: amazon.docx (outdated info)

Answer: "Amazon's stock is currently trading at \$180 per share [1]."

Issue: Price fabricated (not in source documents)

Evaluation:

- Context Relevance: ✓ (1.0 - source retrieved correctly)
- Correctness: ✗ (0.0 - price incorrect)
- Hallucination: ✗ (1.0 - fabricated data)
- Answer Relevance: ✓ (1.0 - addressed question)

Fix: Added web search to fetch current prices

Example 3: Irrelevant Retrieval ✗

Question: "What products does Tesla sell?"

Retrieved Sources: amazon.docx, google.docx (wrong companies)

Answer: "Tesla sells electric vehicles, solar panels, and energy storage products."

Issue: Retrieved wrong documents but answered correctly (from LLM knowledge)

Evaluation:

- Context Relevance: X (0.0 - wrong sources)
- Correctness: ✓ (1.0 - answer correct)
- Hallucination: ~ (0.5 - not grounded in sources)
- Answer Relevance: ✓ (1.0)

Fix: Improved hybrid search parameters, added query expansion

=====

=====

D. REFLECTION

=====

=====

WHAT WORKED

1. Comprehensive Evaluation Framework

Setting up LLM-as-judge with Langfuse from the start enabled data-driven optimization. Running 198 evaluations across 4 metrics gave us confidence in system quality and highlighted specific areas for improvement.

2. Hybrid Search Architecture

Combining BM25 and vector search proved more effective than either alone. The 18% improvement in retrieval quality directly translated to better answer accuracy.

3. Modular Service Design

Separating concerns (news scraping, RAG, sentiment, market data) made the system maintainable and allowed parallel development by team members.

4. Real-time Updates

WebSocket integration for live news and market data created an engaging user experience that keeps information current.

5. Observability from Day One

Langfuse integration provided visibility into every query, LLM call, and retrieval operation, making debugging and optimization straightforward.

WHAT DIDN'T WORK

1. Yahoo Finance Integration

Initially used Yahoo Finance but got blocked due to rate limits. Had to pivot to Nasdaq API and implement CORS proxy in backend.

2. Simple Evaluation Metrics

Early binary pass/fail evaluations weren't nuanced enough. Switched to LLM-as-judge with detailed reasoning for better insights.

3. Static Knowledge Base

RAG system limited to 10 company documents became outdated quickly. Added web search to supplement with current information.

4. Response Latency

Initial 10-second response times hurt UX. Optimized by reducing chunk size, implementing caching, and using streaming responses.

INSIGHTS & FUTURE IMPROVEMENTS

Technical Insights:

- Retrieval is the Bottleneck: Even with advanced LLMs, poor retrieval leads to poor answers. Investing in hybrid search, reranking, and query expansion had the highest ROI.
- Evaluation Requires Multiple Runs: Single LLM-as-judge scores are noisy. We ran 3-5 evaluations per question to get stable metrics.
- Context Matters: Adding document metadata to chunks improved retrieval quality significantly.
- Streaming is Critical: For 3+ second responses, streaming prevents perceived lag.

Future Improvements:

1. Expand Knowledge Base

- Add 100+ company documents
- Integrate 10-K/10-Q SEC filings
- Include earnings call transcripts
- Auto-update documents quarterly

2. Advanced Retrieval

- Implement graph RAG to capture entity relationships
- Add query classification to route to specialized indexes
- Fine-tune embeddings on financial domain
- Implement document timestamping for temporal queries

3. Enhanced Evaluation

- Build golden dataset with 500+ Q&A pairs

- Add human evaluation alongside LLM judge
- Implement A/B testing framework
- Track evaluation metrics over time (regression detection)

4. Performance Optimization

- Add Redis caching for frequent queries
- Implement semantic caching for similar questions
- Batch embedding generation
- Use smaller, faster models for classification tasks

5. User Experience

- Add multi-turn conversation support
- Implement query suggestions
- Add "Explain this chart" visual Q&A
- Create personalized dashboards

6. Production Readiness

- Implement rate limiting and authentication
- Add comprehensive error handling
- Set up monitoring alerts (Grafana)
- Implement blue-green deployment
- Add automated smoke tests in CI/CD

7. Cost Optimization

- Cache embeddings to reduce API calls
- Use GPT-4o-mini instead of GPT-4 where possible
- Implement request batching
- Set up usage monitoring and budget alerts

LESSONS LEARNED

Perspective:

Building this system taught me that EVALUATION IS NOT OPTIONAL—it's the foundation of reliable AI systems. Running 198 evaluations revealed failure modes we never would have caught through manual testing. I also learned that INTEGRATION IS HARDER THAN IMPLEMENTATION—getting RAG, news scraper, frontend, and real-time data to work together smoothly required careful API design and thorough testing.

Team Collaboration:

Clear division of responsibilities (Shehram: backend/evals, Wahab: RAG, Umer: frontend) enabled parallel development but required strong communication and well-defined interfaces.

Production vs. Prototype:

Moving from "works on my machine" to "works in production" required adding observability, error handling, monitoring, and robust API design—easily 50%

of development time.

CONCLUSION

Agentic IRIS successfully combines real-time market data, automated news aggregation, and intelligent RAG-powered Q&A into a cohesive financial intelligence platform. Through rigorous evaluation (198 test runs), we achieved a low hallucination rate (9%) and reasonable retrieval quality (0.45 context relevance), though there's room for improvement in answer correctness (0.39).

KEY ACHIEVEMENTS:

- Production-ready backend with 20+ APIs
- Comprehensive evaluation framework with LLM-as-judge
- RAG system with hybrid search and reranking
- Real-time news and market data integration
- Complete observability with Langfuse

NEXT STEPS:

The system is ready for user testing. Priority improvements include expanding the knowledge base, optimizing performance with caching, and enhancing evaluation coverage.

PROJECT STATISTICS

Technical Metrics:

- Total Backend Endpoints: 20+ REST APIs + 2 WebSocket endpoints
- Total Evaluations Run: 198 (across 4 metrics)
- News Sources Monitored: 10+ companies
- RAG Documents: 10+ company profiles
- Frontend Components: 15+ React components
- Lines of Code: ~15,000+ (Backend: ~8,000, Frontend: ~5,000, RAG: ~2,000)

Performance Metrics:

- API Response Times:
 - * /api/news/all: ~500ms
 - * /api/rag/query-stream: ~3-5s (streaming)
 - * /api/nasdaq/watchlist: ~1-2s
 - * WebSocket latency: <100ms
- System Resources:
 - * Backend: ~500MB RAM
 - * FAISS index: ~50MB
 - * Scraper: ~200MB RAM during execution

Evaluation Results:

- Context Relevance: 0.45 (44% pass rate)
- Correctness: 0.39 (39% pass rate)
- Hallucination: 0.11 (9% hallucination, 91% grounded)
- Answer Relevance: 0.28 (28% pass rate)

=====

TECHNOLOGY STACK

=====

Backend (Shehram):

- FastAPI, Python 3.12, WebSockets
- Playwright for web scraping
- Langfuse for observability
- OpenAI API for sentiment & LLM-as-judge
- Tavily API for web search

RAG System (Wahab):

- FAISS vector store
- OpenAI embeddings (text-embedding-3-small)
- Hybrid search (semantic + BM25)
- GPT-4o-mini for generation
- Langchain for orchestration

Frontend (Umer):

- React 18, TypeScript, Vite
- TailwindCSS
- Recharts for visualizations
- WebSocket client
- Nasdaq API integration

Development Tools:

- Git for version control
- VS Code for development
- Postman for API testing
- Chrome DevTools for debugging
- Langfuse for monitoring

=====

APPENDIX - REPOSITORY STRUCTURE

=====

Agentic_IRIS/

```
├── backend/
│   ├── main.py          # FastAPI app entry point
│   ├── services/        # 5 service modules
│   │   ├── rag_service.py  # RAG integration
│   │   └── news_service.py # News aggregation
```

```

| | | └─ sentiment_service.py # Sentiment analysis
| | | └─ tavily_service.py    # Web search
| | |   └─ stock_service.py   # Market data
| | └─ requirements.txt       # Python dependencies
|
| └─ rag/
| | └─ rag_engine.py          # Core RAG implementation
| | └─ document_loader.py     # Document processing
| | └─ chunker.py             # Semantic chunking
| | └─ embeddings.py          # Embedding generation
| | └─ vector_store.py        # FAISS integration
| | └─ query_processor.py     # Query expansion & reranking
| | └─ langfuse_evals.py      # Evaluation framework
| |   └─ documents/           # 10 company DOCX files
|
| └─ frontend/
| | └─ src/
| | | └─ pages/               # Home, Graph, Chat pages
| | | └─ components/         # UI components
| | |   └─ hooks/            # WebSocket, data fetching
| | └─ package.json           # Node dependencies
|
| └─ scraper/
| | └─ src/
| |   └─ run_scraper.py       # Playwright scraper
|
| └─ CONTRIBUTIONS.md         # Team contributions
| └─ FINAL_REPORT.md          # This report (Markdown)
| └─ FINAL_REPORT_PLAIN.txt    # This report (Plain Text)
| └─ SYSTEM_ARCHITECTURE.md    # Mermaid diagrams
└─ README.md                  # Project overview

```

```

=====
=====
REFERENCES & RESOURCES
=====
=====

```

Documentation:

- Langfuse: <https://langfuse.com/docs>
- FAISS: <https://faiss.ai/>
- OpenAI API: <https://platform.openai.com/docs>
- Playwright: <https://playwright.dev/>
- FastAPI: <https://fastapi.tiangolo.com/>
- React: <https://react.dev/>

Research Papers:

- RAG: Retrieval-Augmented Generation for Knowledge-Intensive NLP Tasks
- Dense Passage Retrieval for Open-Domain Question Answering
- ColBERT: Efficient and Effective Passage Search via Contextualized Late Interaction over BERT

=====

END OF REPORT

=====

This report reflects the state of the system as of December 2024.
For the latest updates, see: <https://github.com/Shehrambaig/iris>

Report Prepared By: Shehram Baig

Team Members: Mr. Wahab (RAG Engine), Umer (Frontend)